

6.824 2018 Lecture 2: Infrastructure: RPC and threads

Most commonly-asked question: Why Go?

6.824 used C++ for many years

C++ worked out well

but students spent time tracking down pointer and alloc/free bugs
and there's no very satisfactory C++ RPC package

Go is a bit better than C++ for us

good support for concurrency (goroutines, channels, &c)

good support for RPC

garbage-collected (no use after freeing problems)

type safe

threads + GC is particularly attractive!

We like programming in Go

relatively simple and traditional

After the tutorial, use https://golang.org/doc/effective_go.html

Russ Cox will give a guest lecture March 8th

Threads

threads are a useful structuring tool

Go calls them goroutines; everyone else calls them threads

they can be tricky

Why threads?

They express concurrency, which shows up naturally in distributed systems

I/O concurrency:

While waiting for a response from another server, process next request

Multicore:

Threads run in parallel on several cores

Thread = "thread of execution"

threads allow one program to (logically) execute many things at once

the threads share memory

each thread includes some per-thread state:

program counter, registers, stack

How many threads in a program?

Sometimes **driven by structure**

e.g. one thread per client, one for background tasks

Sometimes driven by desire for multi-core parallelism

so one active thread per core

the Go runtime automatically schedules runnable goroutines on available cores

Sometimes **driven by desire** for I/O concurrency

the number is determined by latency and capacity

keep increasing until throughput stops growing

Go threads are pretty cheap

100s or 1000s are fine, but maybe not millions

Creating a thread is more expensive than a method call

Threading challenges:

sharing data

one thread reads data that another thread is changing?

e.g. two threads do `count = count + 1`

this is a "race" -- and is usually a bug

-> use Mutexes (or other synchronization)

-> or avoid sharing

coordination between threads

how to wait for all Map threads to finish?

-> use Go channels or WaitGroup

granularity of concurrency

coarse-grained -> simple, but little concurrency/parallelism

fine-grained -> more concurrency, more races and deadlocks

What is a crawler?

goal is to fetch all web pages, e.g. to feed to an indexer

web pages form a graph

multiple links to each page

graph has cycles

Crawler challenges

- Arrange for I/O concurrency

- Fetch many URLs at the same time

- To increase URLs fetched per second

- Since network latency is much more of a limit than network capacity

- Fetch each URL only *once*

- avoid wasting network bandwidth

- be nice to remote servers

- => Need to remember which URLs visited

- Know when finished

Crawler solutions [crawler.go link on schedule page]

Serial crawler:

- the "fetched" map avoids repeats, breaks cycles

- it's a single map, passed by reference to recursive calls

- but: fetches only one page at a time

ConcurrentMutex crawler:

- Creates a thread for each page fetch

- Many concurrent fetches, higher fetch rate

- The threads share the fetched map

- Why the Mutex (== lock)?

- Without the lock:

- Two web pages contain links to the same URL

- Two threads simultaneously fetch those two pages

- T1 checks fetched[url], T2 checks fetched[url]

- Both see that url hasn't been fetched

- Both fetch, which is wrong

- Simultaneous read and write (or write+write) is a "race"

- And often indicates a bug

- The bug may show up only for unlucky thread interleavings

- What will happen if I comment out the Lock()/Unlock() calls?

- go run crawler.go

- go run -race crawler.go

- The lock causes the check and update to be atomic

- How does it decide it is done?

- sync.WaitGroup

- implicitly waits for children to finish recursive fetches

ConcurrentChannel crawler

- a Go channel:

- a channel is an object; there can be many of them

- ch := make(chan int)

- a channel lets one thread send an object to another thread

- ch <- x

- the sender waits until some goroutine receives

- y := <- ch

- for y := range ch

- a receiver waits until some goroutine sends

- so you can use a channel to both communicate and synchronize

- several threads can send and receive on a channel

- remember: **sender blocks until the receiver receives!**

- may be dangerous to hold a lock while sending...

ConcurrentChannel master()

- master() creates a worker goroutine to fetch each page

- worker() sends URLs on a channel

- multiple workers send on the single channel

- master() reads URLs from the channel

- [diagram: master, channel, workers]

- No need to **lock the fetched map, because it isn't shared!**

- Is there any shared data?

- The channel

- The slices and strings sent on the channel

- The arguments master() passes to worker()

When to use sharing and locks, versus channels?

- Most problems can be solved in either style

- What makes the most sense depends on how the programmer thinks

```
state -- sharing and locks
communication -- channels
waiting for events -- channels
```

Use Go's race detector:

```
https://golang.org/doc/articles/race_detector.html
go test -race
```

Remote Procedure Call (RPC)

a key piece of distributed system machinery; all the labs use RPC
goal: easy-to-program client/server communication

RPC message diagram:

```
Client          Server
request--->
<---response
```

RPC tries to mimic local fn call:

```
Client:
  z = fn(x, y)
Server:
  fn(x, y) {
    compute
    return z
  }
```

Rarely this simple in practice...

Software structure

```
client app      handlers
  stubs         dispatcher
RPC lib         RPC lib
  net ----- net
```

Go example: kv.go link on schedule page

A toy key/value storage server -- Put(key,value), Get(key)->value
Uses Go's RPC library

Common:

You have to declare Args and Reply struct for each RPC type

Client:

```
connect()'s Dial() creates a TCP connection to the server
Call() asks the RPC library to perform the call
  you specify server function name, arguments, place to put reply
  library marshalls args, sends request, waits, unmarshalls reply
  return value from Call() indicates whether it got a reply
  usually you'll also have a reply.Err indicating service-level failure
```

Server:

```
Go requires you to declare an object with methods as RPC handlers
You then register that object with the RPC library
You accept TCP connections, give them to RPC library
The RPC library
  reads each request
  creates a new goroutine for this request
  unmarshalls request
  calls the named method (dispatch)
  marshalls reply
  writes reply on TCP connection
The server's Get() and Put() handlers
  Must lock, since RPC library creates per-request goroutines
  read args; modify reply
```

A few details:

Binding: how does client know who to talk to?

For Go's RPC, server name/port is an argument to Dial

Big systems have some kind of name or configuration server

Marshalling: format data into packets

Go's RPC library can pass strings, arrays, objects, maps, &c

Go passes pointers by copying (server can't directly use client pointer)

Cannot pass channels or functions

RPC problem: what to do about failures?

e.g. lost packet, broken network, slow server, crashed server

What does a failure look like to the client RPC library?

Client never sees a response from the server

Client does *not* know if the server saw the request!

Maybe server never saw the request

Maybe server executed, crashed just before sending reply

Maybe server executed, but network died just before delivering reply

[diagram of lost reply]

Simplest failure-handling scheme: "best effort"

Call() waits for response for a while

If none arrives, re-send the request

Do this a few times

Then give up and return an error

Q: is "best effort" easy for applications to cope with?

A particularly bad situation:

client executes

Put("k", 10);

Put("k", 20);

both succeed

what will Get("k") yield?

[diagram, timeout, re-send, original arrives late]

Q: **is best effort ever OK?**

read-only operations

operations that do nothing if repeated

e.g. DB checks if record has already been inserted

Better RPC behavior: "at most once"

idea: **server RPC code detects** duplicate requests

returns previous reply instead of re-running handler

Q: how to detect a duplicate request?

client includes unique ID (XID) with each request

uses same XID for re-send

server:

if seen[xid]:

 r = old[xid]

else

 r = handler()

 old[xid] = r

 seen[xid] = true

some at-most-once complexities

this will come up in lab 3

how to ensure XID is unique?

big random number?

combine unique client ID (ip address?) with sequence #?

server must eventually discard info about old RPCs

when is discard safe?

idea:

each client has a unique ID (perhaps a big random number)

per-client RPC sequence numbers

client includes "seen all replies <= X" with every RPC

much like TCP sequence #s and acks

or only allow client one outstanding RPC at a time

arrival of seq+1 allows server to discard all <= seq

how to handle dup req while original is still executing?

server doesn't know reply yet

idea: "pending" flag per executing RPC; wait or ignore

What if an at-most-once server crashes and re-starts?

if at-most-once duplicate info in memory, server will forget

and accept duplicate requests after re-start

maybe it should write the duplicate info to disk

maybe replica server should also replicate duplicate info

Go RPC is a simple form of "at-most-once"

- open TCP connection

- write request to TCP connection

- Go RPC never re-sends a request

 - So server won't see duplicate requests

- Go RPC code returns an error if it doesn't get a reply

 - perhaps after a timeout (from TCP)

 - perhaps server didn't see request

 - perhaps server processed request but server/net failed before reply came back

What about "exactly once"?

- unbounded retries plus duplicate detection plus fault-tolerant service

Lab 3