

## 6.824 2018 Lecture 3: GFS

## The Google File System

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung  
SOSP 2003

Why are we reading this paper?

- the file system used for map/reduce
- main themes of 6.824 show up in this paper
  - trading consistency for simplicity and performance
  - motivation for subsequent designs

good systems paper -- details from apps all the way to network performance, fault-tolerance, consistency

influential

- many other systems use GFS (e.g., Bigtable, Spanner @ Google)
- HDFS (Hadoop Distributed File System) based on GFS

What is consistency?

A correctness condition

Important but difficult to achieve when data is replicated

especially when application access it concurrently

[diagram: simple example, single machine]

if an application writes, what will a later read observe?

what if the read is from a different application?

but with replication, each write must also happen on other machines

[diagram: two more machines, reads and writes go across]

Clearly we have a problem here.

Weak consistency

read() may return stale data --- not the result of the most recent write

Strong consistency

read() always returns the data from the most recent write()

General tension between these:

strong consistency is easy for application writers

strong consistency is bad for performance

weak consistency has good performance and is easy to scale to many servers

weak consistency is complex to reason about

Many trade-offs give rise to different correctness conditions

These are called "consistency models"

First peek today; will show up in almost every paper we read this term

"Ideal" consistency model

Let's go back to the single-machine case

Would be nice if a replicated FS behaved like a non-replicated file system

[diagram: many clients on the same machine accessing files on single disk]

If one application writes, later reads will observe that write

What if two application concurrently **write to the same file?**

Q: what happens on a single machine?

In file systems often undefined --- file may have some mixed content

What if two application concurrently **write to the same directory**

Q: what happens on a single machine?

One goes first, the other goes second (use locking)

**Challenges to achieving ideal consistency**

Concurrency -- as we just saw; plus there are many disks in reality

Machine failures -- any operation can fail to complete

Network partitions -- may not be able to reach every machine/disk

Why are these challenges difficult to overcome?

Requires communication between clients and servers

May cost performance

Protocols can become complex --- see next week

Difficult to implement system correctly

Many systems in 6.824 don't provide ideal

GFS is one example

GFS goals:

With so many machines, failures are common

must tolerate

assume a machine fails once per year

w/ 1000 machines, ~3 will fail per day.  
 High-performance: many concurrent readers and writers  
 Map/Reduce jobs read and store final result in GFS  
 Note: \*not\* the temporary, intermediate files  
 Use network efficiently: save bandwidth  
 These challenges difficult combine with "ideal" consistency

## High-level design / Reads

[Figure 1 diagram, master + chunkservers]  
 Master stores directories, files, names, open/read/write  
 But not POSIX  
 100s of Linux chunk servers with disks  
 store 64MB chunks (an ordinary Linux file for each chunk)  
 each chunk replicated on three servers  
 Q: Besides availability of data, what does 3x replication give us?  
 load balancing for reads to hot files  
 affinity  
 Q: why not just store one copy of each file on a RAID'd disk?  
 RAID isn't commodity  
 Want fault-tolerance for whole machine; not just storage device  
 Q: why are the chunks so big?  
 amortizes overheads, reduces state size in the master

GFS master server knows directory hierarchy  
 for directory, what files are in it  
 for file, knows chunk servers for each 64 MB  
 master keeps state in memory  
 64 bytes of metadata per each chunk  
 master has private recoverable database for metadata  
 operation log flushed to disk  
 occasional asynchronous compression info checkpoint  
 N.B.: != the application checkpointing in Å§2.7.2  
 master can recovery quickly from power failure  
 shadow masters that lag a little behind master  
 can be promoted to master

## Client read:

send file name and chunk index to master  
 master replies with set of servers that have that chunk  
 response includes version # of chunk  
 clients cache that information  
 ask nearest chunk server  
 checks version #  
 if version # is wrong, re-contact master

## Writes

[Figure 2-style diagram with file offset sequence]  
 Random client write to existing file  
 client asks master for chunk locations + primary  
 master responds with chunk servers, version #, and who is primary  
 primary has (or gets) 60s lease  
 client computes chain of replicas based on network topology  
 client sends data to first replica, which forwards to others  
 pipelines network use, distributes load  
 replicas ack data receipt  
 client tells primary to write  
 primary assign sequence number and writes  
 then tells other replicas to write  
 once all done, ack to client  
 what if there's another concurrent client writing to the same place?  
 client 2 get sequenced after client 1, overwrites data  
 now client 2 writes again, this time gets sequenced first (C1 may be slow)  
 writes, but then client 1 comes and overwrites  
 => all replicas have same data (= consistent), but mix parts from C1/C2  
 (= NOT defined)

## Client append (not record append)

same deal, but may put parts from C1 and C2 in any order  
 consistent, but not defined  
 or, if just one client writes, no problem -- both consistent and defined

## Record append

### Client record append

- client asks master for chunk locations
- client pushes data to replicas, but specifies no offset
- client contacts primary when data is on all chunk servers
  - primary assigns sequence number
  - primary checks if append fits into chunk
    - if not, pad until chunk boundary
  - primary picks offset for append
  - primary applies change locally
  - primary forwards request to replicas
- let's saw R3 fails mid-way through applying the write
  - primary detects error,** tells client to try again
- client retries after contacting master
  - master has perhaps brought up R4 in the meantime (or R3 came back)
  - one replica now has a gap in the byte sequence, so can't just append
  - pad to **next available offset** across all replicas
  - primary and secondaries apply writes
  - primary responds to client after receiving acks from all replicas

### Housekeeping

- Master can appoint new primary if master doesn't refresh lease
- Master replicates chunks if number replicas drop below some number
- Master rebalances replicas

### Failures

- Chunk servers are easy to replace
  - failure may cause some clients to retry (& duplicate records)
- Master: down -> GFS is unavailable
  - shadow master can serve read-only operations, which may return stale data
- Q: Why not write operations?
  - split-brain syndrome (see next lecture)

### Does GFS achieve "ideal" consistency?

**Two cases: directories and files**

Directories: yes, but...

- Yes: strong consistency (only one copy)
- But: master not always available & scalability limit

Files: not always

Mutations with atomic appends

- record can be duplicated at two offsets
- while other replicas may have a hole at one offset

Mutations without atomic append

- data of several clients maybe intermingled
- if you care, use atomic append or a temporary file and atomically rename

An "unlucky" client can read stale data for short period of time

- A failed mutation leaves chunks inconsistent
- The primary chunk server updated chunk
- But then failed and the replicas are out of date
- A client may read an not-up-to-date chunk
- When client refreshes lease it will learn about new version #

Authors claims weak consistency is not a big problems for apps

- Most file updates are **append-only** updates
  - Application can use **UID** in append records to detect duplicates
  - Application **may just read less data** (but not stale data)
- Application can **use temporary files and atomic rename**

### Performance (Figure 3)

- huge aggregate throughput for read (3 copies, striping)
  - 125 MB/sec in aggregate
  - Close to saturating network
- writes to different files lower than possible maximum
  - authors blame their network stack
  - it causes delays in propagating chunks from one replica to next
- concurrent appends file
  - limited by the server that stores last chunk
- numbers and specifics have changed a lot in 15 years!

### Summary

- case study of performance, fault-tolerance, consistency
  - specialized for MapReduce applications

- what works well in GFS?

  - huge sequential reads and writes

  - appends

  - huge throughput (3 copies, striping)

  - fault tolerance of data (3 copies)

- what less well in GFS?

  - fault-tolerance of master

  - small files (master a bottleneck)

  - clients may see stale data

  - appends maybe duplicated

#### References

- <http://queue.acm.org/detail.cfm?id=1594206> (discussion of gfs evolution)

- <http://highscalability.com/blog/2010/9/11/googles-colossus-makes-search-real-time-by-dumping-mapreduce.html>