

6.824 2018 Lecture 4: Primary/Backup Replication

Today

- Primary/Backup Replication for Fault Tolerance
- Case study of VMware FT, an extreme version of the idea

Fault tolerance

- we'd like a service that continues despite failures
- some ideal properties:
 - available: still useable despite [some class of] failures
 - strongly consistent: looks just like a single server to clients
 - transparent to clients
 - transparent to server software
 - efficient

What failures will we try to cope with?

- Fail-stop failures
- Independent failures
- Network drops some/all packets
- Network partition

But not:

- Incorrect execution
- Correlated failures
- Configuration errors
- Malice

Behaviors

- Available (e.g. if one server halts)
- Wait (e.g. if network totally fails)
- Stop forever (e.g. if multiple servers crash)
- Malfunction (e.g. if h/w computes incorrectly, or software has a bug)

Core idea: replication

- *Two* servers (or more)
- Each replica keeps state needed for the service
- If one replica fails, others can continue

Example: fault-tolerant MapReduce master

- lab 1 workers are already fault-tolerant, but not master
- master is a "single point of failure"
- can we have two masters, in case one fails?
- [diagram: M1, M2, workers]
- state:
 - worker list
 - which jobs done
 - which workers idle
 - TCP connection state
 - program memory and stack
 - CPU registers

Big Questions:

- What state to replicate?
- Does primary have to wait for backup?
- When to cut over to backup?
- Are anomalies visible at cut-over?
- How to bring a replacement up to speed?

Two main approaches:

- State transfer
 - "Primary" replica executes the service
 - Primary sends [new] state to backups
- Replicated state machine
 - All replicas execute all operations
 - If same start state,
 - same operations,
 - same order,
 - deterministic,
 - then same end state

State transfer is simpler

But state may be large, slow to transfer
VM-FT uses replicated state machine

Replicated state machine can be more efficient

If operations are small compared to data
But complex to get right
Labs 2/3/4 use replicated state machines

At what level to define a replicated state machine?

K/V put and get?

"application-level" RSM

usually requires server and client modifications

can be efficient; primary only sends high-level operations to backup

x86 instructions?

might allow us to replicate any existing server w/o modification!

but requires much more detailed primary/backup synchronization

and we have to deal with interrupts, DMA, weird x86 instructions

The design of a Practical System for Fault-Tolerant Virtual Machines
Scales, Nelson, and Venkitachalam, SIGOPS OSR Vol 44, No 4, Dec 2010

Very ambitious system:

Goal: fault-tolerance for existing server software

Goal: clients should not notice a failure

Goal: no changes required to client or server software

Very ambitious!

Overview

[diagram: app, O/S, VM-FT underneath, shared disk, network, clients]

words:

hypervisor == monitor == VMM (virtual machine monitor)

app and O/S are "guest" running inside a virtual machine

two machines, primary and backup

shared disk for persistent storage

shared so that bringing up a new backup is faster

primary sends all inputs to backup over logging channel

Why does this idea work?

It's a replicated state machine

Primary and backup boot with same initial state (memory, disk files)

Same instructions, same inputs -> same execution

All else being equal, primary and backup will remain identical

What sources of divergence must we guard against?

Many instructions are guaranteed to execute exactly the same on primary and backup.

As long as memory+registers are identical, which we're assuming by induction.

When might execution on primary differ from backup?

Inputs from external world (the network).

Data read from storage server.

Timing of interrupts.

Instructions that aren't pure functions of state, such as cycle counter.

Races.

Examples of divergence?

They all sound like "if primary fails, clients will see inconsistent story from backup."

Lock server grants lock to client C1, rejects later request from C2.

Primary and backup had better agree on input order!

Otherwise, primary fails, backup now tells clients that C2 holds the lock.

Lock server revokes lock after one minute.

Suppose C1 holds the lock, and the minute is almost exactly up.

C2 requests the lock.

Primary might see C2's request just before timer interrupt, reject.

Backup might see C2's request just after timer interrupt, grant.

So: backup must see same events, in same order, at same point in instruction stream.

Example: timer interrupts

Goal: primary and backup should see interrupt at exactly the same point in execution

i.e. between the same pair of executed instructions

Primary:

FT fields the timer interrupt

FT reads instruction number from CPU

FT sends "timer interrupt at instruction X" on logging channel

FT delivers interrupt to primary, and resumes it

(this relies on special support from CPU to count instructions, interrupt after X)

Backup:

ignores its own timer hardware

FT sees log entry *before* backup gets to instruction X

FT tells CPU to interrupt at instruction X

FT mimics a timer interrupt, resumes backup

Example: disk/network input

Primary and backup *both* ask h/w to read

FT intercepts, ignores on backup, gives to real h/w on primary

Primary:

FT tells the h/w to DMA data into FT's private "bounce buffer"

At some point h/w does DMA, then interrupts

FT gets the interrupt

FT pauses the primary

FT copies the bounce buffer into the primary's memory

FT simulates an interrupt to primary, resumes it

FT sends the data and the instruction # to the backup

Backup:

FT gets data and instruction # from log stream

FT tells CPU to interrupt at instruction X

FT copies the data during interrupt

Why the bounce buffer?

I.e. why wait until primary/backup aren't executing before copying the data?

We want the data to appear in memory at exactly the same point in execution of the primary and backup.

Otherwise they may diverge.

Note that the backup must lag by one event (one log entry)

Suppose primary gets an interrupt, or input, after instruction X

If backup has already executed past X, it cannot handle the input correctly

So backup FT can't start executing at all until it sees the first log entry

Then it executes just to the instruction # in that log entry

And waits for the next log entry before restarting backup

Example: non-functional instructions

even if primary and backup have same memory/registers,

some instructions still execute differently

e.g. reading the current time or cycle count or processor serial #

Primary:

FT sets up the CPU to interrupt if primary executes such an instruction

FT executes the instruction and records the result

sends result and instruction # to backup

Backup:

backup also interrupts when it tries to execute that instruction

FT supplies value that the primary got

What about disk/network output?

Primary and backup both execute instructions for output

Primary's FT actually does the output

Backup's FT discards the output

But: the paper's Output Rule (Section 2.2) says primary must tell backup when it produces output, and delay the output until the backup says it has received the log entry.

Why the Output Rule?

Suppose there was no Output Rule.

The primary emits output immediately.

Suppose the primary has seen inputs I1 I2 I3, then emits output.

The backup has received I1 and I2 on the log.

The primary crashes and the packet for I3 is lost by the network.

Now the backup will go live without having processed I3.

But some client has seen output reflecting the primary having executed I3.

So that client may see anomalous state if it talks to the service again.

So: the primary doesn't emit output until it knows that the backup has seen all inputs up to that output.

The Output Rule is a big deal

Occurs in some form in all replication systems

A serious constraint on performance

An area for application-specific cleverness

Eg. maybe no need for primary to wait before replying to read-only operation

FT has no application-level knowledge, must be conservative

Q: What if the primary crashes just after getting ACK from backup, but before the primary emits the output?

Does this mean that the output won't ever be generated?

A: Here's what happens when the primary fails and the backup takes over.

The backup got some log entries from the primary.

The backup continues executing those log entries WITH OUTPUT SUPPRESSED.

After the last log entry, the backup starts emitting output

In our example, the last log entry is I3

So after input I3, the client will start emitting outputs

And thus it will emit the output that the primary failed to emit

Q: But what if the primary crashed *after* emitting the output?

Will the backup emit the output a *second* time?

A: Yes.

OK for TCP, since receivers ignore duplicate sequence numbers.

OK for writes to shared disk, since backup will write same data to same block #.

Duplicate output at cut-over is pretty common in replication systems

Not always possible for clients to ignore duplicates

For example, if output is vending money from an ATM machine

Q: Does FT cope with network partition -- could it suffer from split brain?

E.g. if primary and backup both think the other is down.

Will they both "go live"?

A: The shared disk breaks the tie.

Shared disk server supports atomic test-and-set.

Only one of primary/backup can successfully test-and-set.

If only one is alive, it will win test-and-set and go live.

If both try, one will lose, and halt.

Shared storage is single point of failure

If shared storage is down, service is down

Maybe they have in mind a replicated storage system

Q: Why don't they support multi-core?

Performance (table 1)

FT/Non-FT: impressive!

little slow down

Logging bandwidth

Directly reflects disk read rate + network input rate

18 Mbit/s for my-sql

These numbers seem low to me

Applications can read a disk at at least 400 megabits/second

So their applications aren't very disk-intensive

When might FT be attractive?

Critical but low-intensity services, e.g. name server.

Services whose software is not convenient to modify.

What about replication for high-throughput services?

People use application-level replicated state machines for e.g. databases.

The state is just the DB, not all of memory+disk.

The events are DB commands (put or get), not packets and interrupts.
Result: less fine-grained synchronization, less overhead.
GFS use application-level replication, as do Lab 2 &c

Summary:

Primary-backup replication

VM-FT: clean example

How to cope with partition without single point of failure?

Next lecture

How to get better performance?

Application-level replicated state machines

VMware KB (#1013428) talks about multi-CPU support. VM-FT may have switched from a replicated state machine approach to the state transfer approach, but unclear whether that is true or not.

<http://www.wooditwork.com/2014/08/26/whats-new-vsphere-6-0-fault-tolerance/>

<http://www.tomsitpro.com/articles/vmware-vsphere-6-fault-tolerance-multi-cpu,1-2439.html>

<https://labs.vmware.com/academic/publications/retrace>